


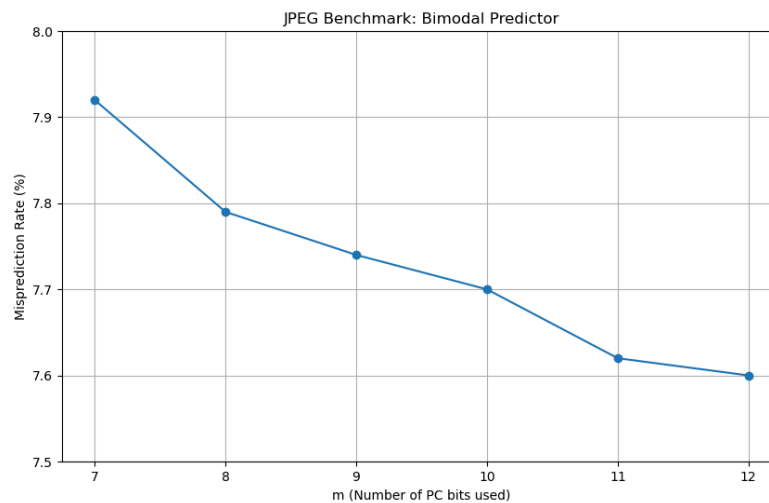
Name: Varun M
ROLL No: EE20B149

1. Bimodal Predictor

- GCC



- JPEG



Analysis:

GCC Benchmark trend:

- There's a significant improvement in prediction accuracy as m increases, particularly in the lower m values.
- The rate of improvement is higher between m=7 and m=10, with each step providing substantial gains.

- The improvement starts to slow down after $m=10$, suggesting diminishing returns for larger predictor tables.
- Even at $m=12$, the misprediction rate is still relatively high at 12.47%, indicating that the GCC benchmark has complex branch behavior that remains challenging to predict accurately.

JPEG Benchmark Trends:

- The overall misprediction rates are much lower compared to GCC, starting at 7.92% for $m=7$.
- There's only a slight improvement as m increases, with the total reduction being just 0.32 percentage points from $m=7$ to $m=12$.
- The rate of improvement is relatively consistent across all m values, without any sharp drops or plateaus.

Similarities between benchmarks:

- Both benchmarks show an overall downward trend in misprediction rates as m increases.
- For both benchmarks, the improvements become less pronounced at higher m values, indicating diminishing returns for larger predictor tables.

Differences:

- Initial misprediction rates: GCC starts with a much higher misprediction rate (26.65% at $m=7$) compared to JPEG (7.92% at $m=7$).
- Magnitude of improvement: GCC shows a dramatic improvement, reducing misprediction rate by 14.18 percentage points, while JPEG only improves by 0.32 percentage points.
- Sensitivity to m : GCC's performance is much more sensitive to changes in m , especially for lower m values. JPEG's performance, on the other hand, is only marginally affected by increases in m .
- Final misprediction rates: Even at $m=12$, GCC's misprediction rate (12.47%) is still significantly higher than JPEG's worst rate (7.92% at $m=7$).

Overall Observation:

- The GCC benchmark, which represents a compiler workload, seems to have more complex and harder-to-predict branch behavior. This could be due to the diverse code patterns and decision structures typically found in compiler operations.

- The JPEG benchmark, representing image processing, appears to have more predictable branch behavior. This might be due to more regular, loop-based operations common in image processing algorithms.
- For GCC, increasing the predictor size (m) provides substantial benefits, especially up to $m=10$. This suggests that for workloads similar to compilers, larger branch predictors can significantly improve performance.
- For JPEG, the benefit of increasing predictor size is minimal. This implies that for workloads similar to image processing, even smaller predictors ($m=7$ or 8) might be sufficient, and allocating more resources to the branch predictor may not be cost-effective.
- If designing a general-purpose processor that needs to handle both types of workloads efficiently, a compromise might be to choose $m=10$ or $m=11$. This would provide most of the benefit for GCC-like workloads while not oversizing the predictor for JPEG-like workloads.

Proposed Branch Predictor design:

The maximum allowed $m = 16$ as the storage for bimodal is 2^m . 2 bits
 $2^{m+1} \leq 2^{17}$

$m \leq 16$

For Gcc benchmark

As m increases, the misprediction rate decreases significantly, especially from $m = 7$ to $m = 11$.

or the **gcc benchmark**, the misprediction rates for different values of m are:

- $m = 10$: 15.67%
- $m = 11$: 13.65%
- $m = 12$: 12.47%

While the **misprediction rate** at $m=11$ is **13.65%**, which is slightly higher than $m=12$ (12.47%), it represents a significant drop compared to $m=10$ (15.67%). The decrease in misprediction rate from $m=10$ to $m=11$ is **2.02%**, while the further improvement from $m=11$ to $m=12$ is only **1.18%**.

This indicates that $m = 11$ offers a substantial reduction in misprediction rate compared to $m = 10$, while not requiring the additional storage that comes with $m = 12$.

2. Storage Cost

For **m = 11**, the storage requirement is:

$2^{11} \times 2/8 = 512$ bytes This is a reasonable storage requirement, well within the **16kB** budget. Compared to **m = 10** (which uses **256 bytes**), it doubles the storage requirement but provides a meaningful improvement in prediction accuracy.

3. Trade-off Between Accuracy and Cost

The trade-off between storage and accuracy suggests that **m = 11** strikes a balance:

- **m = 10** offers lower storage but a higher misprediction rate.
- **m = 12** has a slightly better misprediction rate but comes at the cost of larger storage (1024 bytes).

Thus, **m = 11** can be considered the **sweet spot**, offering a **significant reduction in misprediction rate** over **m = 10** while avoiding the diminishing returns of increasing to **m = 12**.

Proposed Design for GCC with m = 11:

- **Misprediction rate: 13.65%.**
- **Storage requirement: 512 bytes.**

This configuration offers a good balance between prediction accuracy and storage cost, making it a solid choice for the **gcc benchmark**.

Jpeg Benchmark:

The misprediction rates for the **jpeg** benchmark for various values of mmm are:

- **m = 7: 7.92%**
- **m = 8: 7.79%**
- **m = 9: 7.74%**
- **m = 10: 7.7%**
- **m = 11: 7.62%**
- **m = 12: 7.6%**

Analysis:

- The misprediction rate shows very little improvement beyond **m = 8**. The change from **m = 8** to **m = 12** is only **0.19%**.

- Since the improvement in prediction accuracy is minimal after **m = 8**, there's little justification for increasing mmm further.

Storage Consideration:

- For **m = 8**, the storage requirement is: $28 \times 2/8 = 64$ bytes 2^8
- For **m = 9**, the storage requirement is: $2^9 \times 2/8 = 128$ bytes
- For **m = 10**, the storage requirement is: $2^{10} \times 2/8 = 256$ bytes

The predictor for **m = 8** is highly efficient, requiring only **64 bytes**, and further increasing mmm gives marginal improvements in misprediction rate.

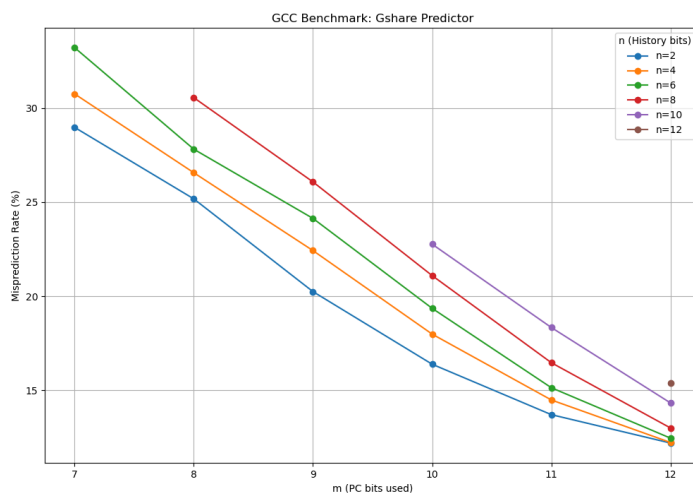
Proposed Design for JPEG:

- **m = 8** is optimal for the **jpeg** benchmark:
 - **Misprediction rate:** 7.79%
 - **Storage requirement:** **64 bytes**.

This configuration provides good prediction accuracy at a minimal storage cost, which leaves ample room within the 16kB budget.

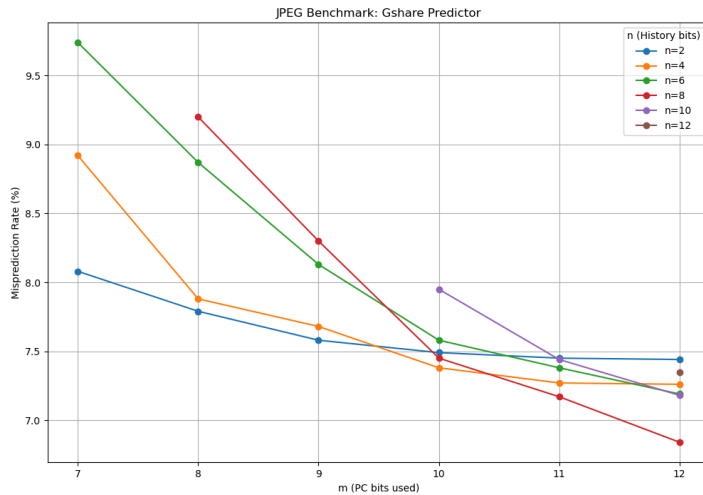
2. Gshare Predictor

GCC



- As m increases for all values of n, the misprediction rate decreases.
- For small values of n (e.g., n=2,4), the misprediction rate starts high but drops significantly as m increases.

- Larger values of n (e.g., $n=10,12$) show a steeper decrease in misprediction rates for smaller values of m , but the benefit diminishes beyond $m=10$.
- The lowest misprediction rates are observed at high m and high n .
- **JPEG**



- For the **jpeg** benchmark, the misprediction rate is consistently lower than that of the gcc benchmark.
- The trend is similar: increasing m reduces the misprediction rate for each value of n .
- The rates start to stabilize at higher values of m and show diminishing returns in prediction accuracy, especially for $n=10$ and $n=12$.
- The misprediction rates for jpeg are relatively stable compared to gcc, with less dramatic improvement as m increases.

ANALYSIS

Trends:

- For **gcc**, larger values of n yield better prediction accuracy, with a significant reduction in misprediction rates as m increases. However, the improvements taper off for values of m greater than 10, where the misprediction rate stabilizes.
- For **jpeg**, the misprediction rates are consistently lower across all configurations, with diminishing improvements as m increases, especially for larger values of n .

Similarities:

- Both benchmarks show a decrease in misprediction rates with increasing m across all n values.
- Both experience diminishing returns in prediction accuracy as m approaches 12.

Differences:

- The **gcc** benchmark shows more variability and a sharper decline in misprediction rates, particularly for smaller values of n .
- The **jpeg** benchmark has more stable misprediction rates and a generally lower misprediction rate overall compared to gcc.

Proposed Architecture:

For the **gcc_trace.txt** benchmark, the misprediction rate drops significantly as m increases. For example, when $n=2$, the misprediction rate decreases from **28.98%** at $m=7$ to **12.2%** at $m=12$. However, after $m=10$, the decrease becomes less notable. Similarly, increasing n offers diminishing returns beyond $n=4$, with the misprediction rate stabilizing around **12.46%** at $m=12$ for $n=6$. Therefore, the optimal configuration for **gcc_trace.txt** is $m=11$ and $n=4$, providing a good balance between performance and storage, as further increases in m and n offer minimal improvement.

For the **jpeg_trace.txt** benchmark, the misprediction rate is less affected by changes in m and n . Starting at **8.08%** for $n=2$, it decreases slightly to **7.44%** at $m=12$. Beyond $m=10$, improvements are minimal, and increasing n has little impact on the misprediction rate. Thus, the best configuration for **jpeg_trace.txt** is $m=10$ and $n=2$, which optimizes performance without excessive storage or power overhead.

Both configurations fit comfortably within the storage limit. For **gcc_trace.txt**, $m=11$ requires **4,096 bits** for the predictor table, while **jpeg_trace.txt** requires **2,048 bits** at $m=10$. These values are well within the available **16kB (131,072 bits)** of storage.

In summary, the optimal configuration for **gcc_trace.txt** is $m=11$ and $n=4$, while for **jpeg_trace.txt**, it is $m=10$ and $n=2$. Both designs effectively minimize misprediction rates while staying within the 16kB storage constraint.